

pa2-actor-critic-acrobot

March 29, 2023

```
[ ]: import numpy as np
import random
import torch
import torch.nn as nn
import torch.nn.functional as F
from collections import namedtuple, deque
import torch.optim as optim
import datetime
import gym
from gym.wrappers.record_video import RecordVideo
import glob
import io
import base64
import matplotlib.pyplot as plt
from IPython.display import HTML
from pyvirtualdisplay import Display
import tensorflow as tf
from IPython import display as ipythondisplay
from PIL import Image
import tensorflow_probability as tfp
```

```
[ ]: class ActorCriticModel(tf.keras.Model):
    """
    Defining policy and value networkss
    """
    def __init__(self,
        ↪action_size, num_layers=2, num_neurons_each_layer=[1024, 512]):
        super(ActorCriticModel, self).__init__()
        self.num_layers = num_layers
        self.neurons_each_layer = num_neurons_each_layer
        self.linears = [tf.keras.layers.Dense(self.neurons_each_layer[0],
        ↪activation=tf.nn.relu)]
        self.linears.extend([tf.keras.layers.Dense(self.neurons_each_layer[i],
        ↪activation=tf.nn.relu) for i in range(1, self.num_layers)])

        #Output Layer for policy
```

```

        self.pi_out = tf.keras.layers.Dense(action_size, activation=tf.nn.
↪softmax)
        #Output Layer for state-value
        self.v_out = tf.keras.layers.Dense(1)

    def call(self, state):
        """
        Computes policy distribution and state-value for a given state
        """
        h = None
        for i in range(self.num_layers+1):
            if i == 0:
                h = tf.nn.relu(self.linears[0](state))
            elif i < self.num_layers:
                h = tf.nn.relu(self.linears[i](h))
            else:
                return self.pi_out(h), self.v_out(h)

```

```

[ ]: class Agent:

    def __init__(self, action_size, num_layers, num_neurons_each_layer, lr=0.
↪001, gamma=0.99, seed = 85):
        self.gamma = gamma
        self.ac_model = ↵
↪ActorCriticModel(action_size=action_size,num_layers=num_layers,num_neurons_each_layer=num_n
        self.ac_model.compile(tf.keras.optimizers.Adam(learning_rate=lr))
        np.random.seed(seed)

    def sample_action(self, state):

        pi,_ = self.ac_model(state)

        action_probabilities = tfp.distributions.Categorical(probs=pi)
        sample = action_probabilities.sample()

        return int(sample.numpy()[0])

    def actor_loss(self, action, pi, delta):

        return -tf.math.log(pi[0,action]) * delta

    def critic_loss(self,delta):
        return delta**2

    @tf.function
    def learn(self, state, action, reward, next_state, done):

```

```

with tf.GradientTape(persistent=True) as tape:
    pi, V_s = self.ac_model(state)
    _, V_s_next = self.ac_model(next_state)

    V_s = tf.squeeze(V_s)
    V_s_next = tf.squeeze(V_s_next)

    delta = reward+((self.gamma)*V_s_next) - V_s
    loss_a = self.actor_loss(action, pi, delta)
    loss_c = self.critic_loss(delta)
    loss_total = loss_a + loss_c

    gradient = tape.gradient(loss_total, self.ac_model.trainable_variables)
    self.ac_model.optimizer.apply_gradients(zip(gradient, self.ac_model.
↪trainable_variables))

```

```

[ ]: import warnings
warnings.filterwarnings("ignore", category=FutureWarning)

def Actor_critic_one_step_return(num_layers, num_neurons_each_layer,
↪LR,num_episodes,seed,TRUNCATION):
    env = gym.make('Acrobot-v1',max_episode_steps=TRUNCATION)

    #Initializing Agent
    agent = Agent( action_size=env.action_space.n, num_layers=num_layers,
↪num_neurons_each_layer=num_neurons_each_layer, lr=LR,seed = seed)
    #Number of episodes
    episodes = num_episodes
    tf.compat.v1.reset_default_graph()

    reward_list = []
    steps = []
    ep_solved = num_episodes

    begin_time = datetime.datetime.now()

    for ep in range(1, episodes + 1):
        state = env.reset().reshape(1,-1)
        done = False
        ep_rew =0
        T =0
        while not done:
            T += 1
            action = agent.sample_action(state) # Action taken in an episode
            next_state, reward, done, info = env.step(action) # Observing the
↪state transitions

```

```

        next_state = next_state.reshape(1,-1)
        ep_rew += reward # Accumulate the reward
        agent.learn(state, action, reward, next_state, done) # Modify the
        ↪weights of neural network based on observed weights
        state = next_state
        reward_list.append(ep_rew)
        steps.append(T)

    if ep % 10 == 0:
        avg_rew = np.mean(reward_list[-10:])
        print('Episode ', ep, 'Reward %f' % ep_rew, 'Average Reward %f' %
        ↪avg_rew)
        if ep>1000 and (np.std(reward_list[-500:])*100/np.mean(reward_list[-500:
        ↪])) < 5:
            ep_solved = num_episodes
            break
    if ep % 100:
        avg_100 = np.mean(reward_list[-100:])
        if avg_100 > -100:
            print('Stopped at Episode ',ep-100)
            ep_solved= ep-100
            break

    time_taken = datetime.datetime.now() - begin_time
    print(time_taken)
    return [reward_list,steps,ep_solved]

```

```

[ ]: SEED = 101
def calculate_mean(array):
    lens = [len(i) for i in array]
    arr = np.ma.empty((np.max(lens),len(array)))
    arr.mask = True
    for idx, l in enumerate(array):
        arr[:len(l),idx] = l
    return arr.mean(axis = -1), arr.std(axis=-1)

```

```

[ ]: import numpy as np
rewards_10, steps_10,ep_10 = [],[],[]
for i in range(3):
    rewards_run, steps_run,ep = Actor_critic_one_step_return(num_layers=2,
    ↪num_neurons_each_layer=[128,128], LR=1e-4,num_episodes=500,seed=
    ↪SEED,TRUNCATION = 300)
    rewards_10.append(rewards_run)
    steps_10.append(steps_run)
    ep_10.append(ep)

reward_10_runs,error_rewards =calculate_mean(rewards_10)

```

```

steps_10_runs,error_steps = calculate_mean(steps_10)
avg_ep = np.mean(ep_10)
print(rewards_10, steps_10,ep_10)

```

```

Episode 10 Reward -300.000000 Average Reward -289.900000
Episode 20 Reward -300.000000 Average Reward -247.000000
Episode 30 Reward -232.000000 Average Reward -282.700000
Episode 40 Reward -300.000000 Average Reward -269.000000
Episode 50 Reward -273.000000 Average Reward -253.500000
Episode 60 Reward -300.000000 Average Reward -282.500000
Episode 70 Reward -300.000000 Average Reward -299.200000
Episode 80 Reward -219.000000 Average Reward -291.900000
Episode 90 Reward -289.000000 Average Reward -232.100000
Episode 100 Reward -274.000000 Average Reward -212.600000
Episode 110 Reward -177.000000 Average Reward -231.900000
Episode 120 Reward -300.000000 Average Reward -276.600000
Episode 130 Reward -300.000000 Average Reward -300.000000
Episode 140 Reward -300.000000 Average Reward -300.000000
Episode 150 Reward -300.000000 Average Reward -300.000000
Episode 160 Reward -300.000000 Average Reward -300.000000
Episode 170 Reward -300.000000 Average Reward -300.000000
Episode 180 Reward -300.000000 Average Reward -300.000000
Episode 190 Reward -300.000000 Average Reward -300.000000
Episode 200 Reward -300.000000 Average Reward -300.000000
Episode 210 Reward -300.000000 Average Reward -300.000000
Episode 220 Reward -300.000000 Average Reward -300.000000
Episode 230 Reward -300.000000 Average Reward -300.000000
Episode 240 Reward -300.000000 Average Reward -300.000000
Episode 250 Reward -300.000000 Average Reward -300.000000
Episode 260 Reward -300.000000 Average Reward -300.000000
Episode 270 Reward -300.000000 Average Reward -300.000000
Episode 280 Reward -300.000000 Average Reward -300.000000
Episode 290 Reward -300.000000 Average Reward -300.000000
Episode 300 Reward -300.000000 Average Reward -300.000000
Episode 310 Reward -300.000000 Average Reward -300.000000
Episode 320 Reward -300.000000 Average Reward -300.000000
Episode 330 Reward -300.000000 Average Reward -300.000000
Episode 340 Reward -300.000000 Average Reward -300.000000
Episode 350 Reward -300.000000 Average Reward -300.000000
Episode 360 Reward -300.000000 Average Reward -300.000000
Episode 370 Reward -300.000000 Average Reward -300.000000
Episode 380 Reward -300.000000 Average Reward -300.000000
Episode 390 Reward -300.000000 Average Reward -300.000000
Episode 400 Reward -300.000000 Average Reward -300.000000
Episode 410 Reward -300.000000 Average Reward -300.000000
Episode 420 Reward -300.000000 Average Reward -300.000000
Episode 430 Reward -300.000000 Average Reward -300.000000

```

[illegible]

Episode 410 Reward -300.000000 Average Reward -300.000000
Episode 420 Reward -300.000000 Average Reward -300.000000
Episode 430 Reward -300.000000 Average Reward -300.000000
Episode 440 Reward -300.000000 Average Reward -300.000000
Episode 450 Reward -300.000000 Average Reward -300.000000
Episode 460 Reward -300.000000 Average Reward -300.000000
Episode 470 Reward -300.000000 Average Reward -300.000000
Episode 480 Reward -300.000000 Average Reward -300.000000
Episode 490 Reward -300.000000 Average Reward -300.000000
Episode 500 Reward -300.000000 Average Reward -300.000000

0:12:03.668842

Episode 10 Reward -300.000000 Average Reward -279.500000
Episode 20 Reward -300.000000 Average Reward -276.300000
Episode 30 Reward -247.000000 Average Reward -257.200000
Episode 40 Reward -300.000000 Average Reward -271.500000
Episode 50 Reward -217.000000 Average Reward -239.700000
Episode 60 Reward -300.000000 Average Reward -285.300000
Episode 70 Reward -295.000000 Average Reward -299.500000
Episode 80 Reward -241.000000 Average Reward -241.000000
Episode 90 Reward -202.000000 Average Reward -191.500000
Episode 100 Reward -154.000000 Average Reward -216.100000
Episode 110 Reward -143.000000 Average Reward -165.200000
Episode 120 Reward -137.000000 Average Reward -171.400000
Episode 130 Reward -103.000000 Average Reward -126.900000
Episode 140 Reward -236.000000 Average Reward -140.200000
Episode 150 Reward -142.000000 Average Reward -150.600000
Episode 160 Reward -95.000000 Average Reward -108.200000
Episode 170 Reward -89.000000 Average Reward -98.900000
Episode 180 Reward -97.000000 Average Reward -115.200000
Episode 190 Reward -103.000000 Average Reward -93.100000
Episode 200 Reward -103.000000 Average Reward -106.800000
Episode 210 Reward -95.000000 Average Reward -105.600000
Episode 220 Reward -73.000000 Average Reward -88.700000
Episode 230 Reward -80.000000 Average Reward -122.900000
Episode 240 Reward -77.000000 Average Reward -110.400000
Episode 250 Reward -94.000000 Average Reward -90.800000
Episode 260 Reward -89.000000 Average Reward -94.800000
Episode 270 Reward -125.000000 Average Reward -90.600000
Episode 280 Reward -77.000000 Average Reward -103.100000
Episode 290 Reward -98.000000 Average Reward -96.600000

Stopped at Episode 195

0:03:59.341325

[[-300.0, -300.0, -300.0, -282.0, -300.0, -300.0, -217.0, -300.0, -300.0,
-300.0, -275.0, -167.0, -226.0, -281.0, -257.0, -286.0, -286.0, -246.0, -146.0,
-300.0, -260.0, -300.0, -300.0, -300.0, -297.0, -300.0, -300.0, -300.0, -238.0,
-232.0, -300.0, -300.0, -262.0, -218.0, -300.0, -216.0, -235.0, -300.0, -259.0,
-300.0, -218.0, -211.0, -300.0, -241.0, -235.0, -157.0, -300.0, -300.0, -300.0,
-273.0, -211.0, -300.0, -300.0, -300.0, -214.0, -300.0, -300.0, -300.0, -300.0,

[illegible]

[illegible]

[illegible]

[illegible]

```

231, 134, 120, 165, 196, 201, 136, 130, 205, 144, 144, 208, 149, 157, 182, 171,
175, 200, 200, 138, 138, 131, 128, 141, 146, 147, 101, 144, 99, 104, 106, 128,
129, 119, 129, 113, 125, 145, 181, 237, 263, 111, 200, 165, 132, 143, 138, 118,
103, 143, 104, 95, 132, 136, 78, 146, 100, 104, 101, 96, 99, 86, 120, 96, 94,
129, 110, 100, 75, 90, 94, 150, 102, 127, 102, 131, 124, 116, 118, 98, 102, 80,
93, 103, 71, 96, 91, 89, 112, 104, 82, 104, 148, 94, 130, 144, 75, 117, 80, 104,
116, 117, 100, 160, 111, 84, 85, 103, 94, 96, 125, 84, 76, 107, 91, 82, 88, 88,
82, 74, 100, 101, 194, 151, 93, 192, 68, 91, 168, 81, 86, 246, 113, 80, 82, 81,
86, 101, 161, 78, 93, 74, 80, 99, 91, 106, 114, 101, 65, 95, 90, 91, 81, 111,
83, 86, 99, 101, 126, 90, 84, 66, 107, 92, 104, 92, 84, 80, 81, 126, 93, 78,
102, 109, 137, 121, 105, 108, 110, 78, 109, 86, 174, 88, 102, 86, 79, 78, 75,
99, 80, 85, 81, 80, 92]] [500, 500, 195]

```

```


[ ]: reward_10_runs,error_rewards =calculate_mean(rewards_10)
steps_10_runs,error_steps = calculate_mean(steps_10)

```


```

[ ]: import matplotlib.pyplot as plt
import numpy as np
plt.figure(figsize = (15,5))
plt.subplot(131)
plt.plot(np.arange(len(reward_10_runs[: -300])), reward_10_runs[: -300], color =_
↳'darkblue')
plt.title("Episodes vs rewards")
plt.xlabel("Episodes")
plt.ylabel('Reward')
plt.subplot(132)
plt.plot(np.arange(len(steps_10_runs[: -300])), steps_10_runs[: -300], color =_
↳'darkblue')
plt.title("Episodes vs steps")
plt.xlabel("Episodes")
plt.ylabel('steps')
plt.subplot(133)
plt.plot(np.arange(len(error_rewards[: -300])), error_rewards[: -300], color =_
↳'darkblue')
plt.title("Episodes vs variance")
plt.xlabel("Episodes")
plt.ylabel('Deviation observed')
plt.show()

```



```
[ ]: import matplotlib.pyplot as plt
import numpy as np
plt.figure(figsize = (15,5))
plt.subplot(131)
plt.plot(np.arange(len(reward_10_runs)), reward_10_runs, color = 'darkblue')
plt.title("Episodes vs rewards")
plt.xlabel("Episodes")
plt.ylabel('Reward')
plt.subplot(132)
plt.plot(np.arange(len(steps_10_runs)), steps_10_runs, color = 'darkblue')
plt.title("Episodes vs steps")
plt.xlabel("Episodes")
plt.ylabel('steps')
plt.subplot(133)
plt.plot(np.arange(len(error_rewards)), error_rewards, color = 'darkblue')
plt.title("Episodes vs variance")
plt.xlabel("Episodes")
plt.ylabel('Deviation observed')
plt.show()
```



```

[ ]: reward = [-308.8,
-204.6,
-250.6,
-330.7,
-377.,0
-284.3,
-192.3,
-186.3,
-131.6,
-136.5,
-158.6,
-165.,0
-165.1,
-159.8,
-133.8,
-162.3,
-173.9,
-172.5,
-154.2,
-136.1,
-152.3,
-154.5,
-127.5,
-137.1,
-219.5,
-113.4,
-122.1,
-122.7,
-114.5,
-106.8,
-126.5,
-111.7,
-111.6,
-92.9,
-103.,0
-97.5,
-87.1,
-105.6,
-91.9,
-94.9]

reward_2 = [-256.000000
-362.200000
-500.000000
-376.100000

```

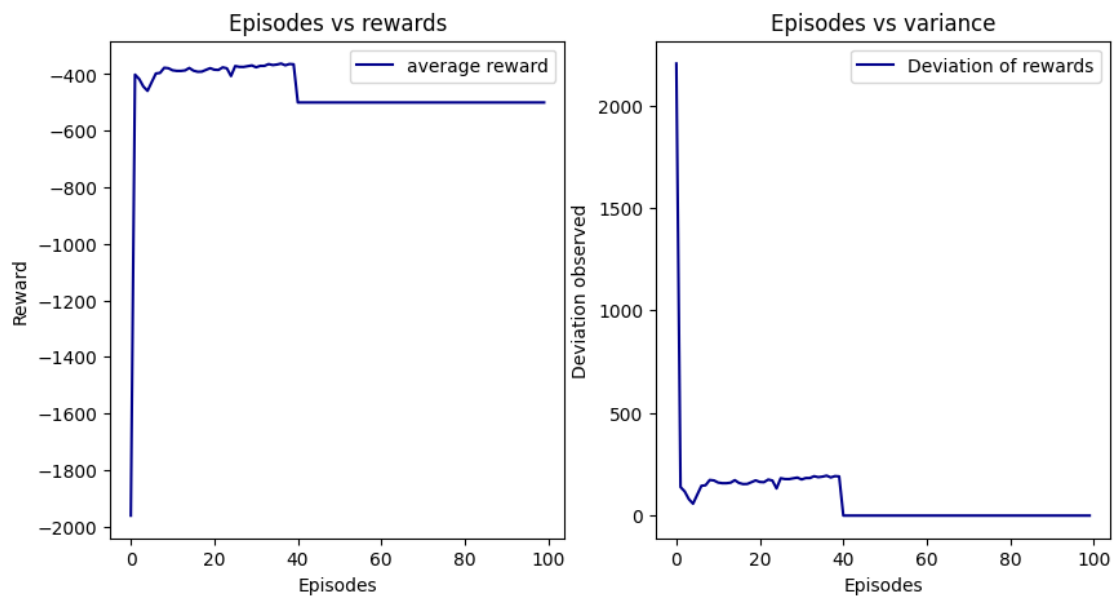
```

-455.100000
-456.100000
-344.200000
-260.300000
  -243.400000
  -240.500000
  -245.800000
  -409.900000
  -435.500000
  -490.500000]
counter = 160
while counter<=1000:
    reward_2.append(-500)
    counter+=10

reward_3 = []
co = 10
while(co <= 1000):
    reward_3.append(-500)
    co += 10

reward = [reward,reward_2, reward_3]

```



```

[ ]: import wandb
      !wandb login --relogin

```

```

import pprint
import numpy as np

sweep_config = {
    'method': 'bayes'
}

metric = {
    'name' : 'num_episodes_tosolve_cartpole_full_ac',
    'goal' : 'minimize'
}
sweep_config['metric'] = metric

parameters_dict = {
    'NUM_NEURONS_EACH_LAYER' : {
        'values': [64,90,128]
    },
    'NUM_LAYERS' : {
        'values': [2,4,8]
    },
    'NUM_EPISODES': {
        'values': [1800,2000,2500]
    }
}

sweep_config['parameters'] = parameters_dict

import math
parameters_dict.update({
    'LR': {
        'distribution': 'uniform',
        'min': (5e-5),
        'max': (5e-4)
    }
})

pprint.pprint(sweep_config)

sweep_id = wandb.sweep(sweep_config, project="Hyper parameter tuning RL full_
↳return - v1")

def train(config = None):
    with wandb.init(config = config):

```



```

config = wandb.config
LR = config.LR                                ## learning rate
NUM_NEURONS_EACH_LAYER = config.NUM_NEURONS_EACH_LAYER ##number of neurons
↳in each hidden layer
NUM_LAYERS = config.NUM_LAYERS                ##number of layers for the
↳neural network
NUM_EPISODES = config.NUM_EPISODES
steps_over_10_runs = []
rewards_over_10_runs = []
episodes = []
min_len = np.inf
for run in range(3):
    print("Run: ",run+1,"
↳-----")

    env = gym.make('CartPole-v1')

    seed = np.random.randint(low=0, high = 100)
    env.seed(seed)
    state_shape = env.observation_space.shape[0]
    action_shape = env.action_space.n

    begin_time = datetime.datetime.now()    #(num_layers,
↳num_neurons_each_layer, LR,num_episodes,seed):
    rewards_each_ep, steps_each_ep, episodes_solved =
↳Actor_critic_one_step_return(NUM_LAYERS,
↳[NUM_NEURONS_EACH_LAYER]*NUM_LAYERS, LR,NUM_EPISODES,seed)
    steps_over_10_runs.append(steps_each_ep)
    rewards_over_10_runs.append(rewards_each_ep)
    episodes.append(episodes_solved)

    if len(rewards_each_ep) < min_len:
        min_len = len(rewards_each_ep)
    time_taken = datetime.datetime.now() - begin_time

    print(time_taken)

rewards = np.array(rewards_over_10_runs)
avg_episodes = int(np.array(episodes).mean())
steps = np.array(steps_over_10_runs)

y, error_reward = calculate_mean(rewards[:avg_episodes])
plt.figure(figsize = (15,5))
plt.subplot(121)
# for i in range(len(rewards_over_10_runs)):
#     plt.plot(np.
↳arange(len(rewards_over_10_runs[i]),rewards_over_10_runs[i],label='run
↳'+str(i))

```

```

plt.plot(np.arange(len(y))+1, y, color='darkblue',label='average')
plt.title("Episodes vs rewards")
plt.xlabel("Number of episodes")
plt.ylabel("Reward per episode")
plt.legend()
plt.subplot(122)
y, error = calculate_mean(steps[:avg_episodes])
plt.plot(np.arange(len(y))+1, y, color='darkblue',label='average')
plt.title("Episodes vs steps ")
plt.xlabel("Number of episodes")
plt.ylabel("Steps taken per episode")
plt.legend()
plt.subplot(133)
plt.plot(np.arange(len(error_reward))+1, error_reward ,
color='darkblue',label= 'Deviation of rewards') #plotting the variance = std
squared
plt.title("Episodes vs Variance ")
plt.xlabel("Number of episodes")
plt.ylabel("Observed Variance")
plt.legend()
plt.savefig('hyp_cartpole_ac_one_step_return'+str(NUM_EPISODES)+str(LR)+'
.jpg', dpi = 250)
plt.show()

for i in range(avg_episodes):
    print("Episode: ",i+1, 'Number of steps: ',round(y[i]))
wandb.log({"num_episodes_tosolve_cartpole_full_ac":avg_episodes})

```

wandb: ERROR Find detailed error logs at: c:\Users\CFI
Workstation\Documents\RL\wandb\debug-cli.CFI Workstation.log
Error: api_key not configured (no-tty). call wandb login [your_api_key]

```

{'method': 'bayes',
 'metric': {'goal': 'minimize',
            'name': 'num_episodes_tosolve_cartpole_full_ac'},
 'parameters': {'LR': {'distribution': 'uniform', 'max': 0.0005, 'min': 5e-05},
                 'NUM_EPISODES': {'values': [1800, 2000, 2500]},
                 'NUM_LAYERS': {'values': [2, 4, 8]},
                 'NUM_NEURONS_EACH_LAYER': {'values': [64, 90, 128]}}}

```

C:\Users\CFI Workstation\AppData\Local\Packages\PythonSoftwareFoundation.Python.
3.10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\notebook\utils.py:280: DeprecationWarning: distutils Version classes
are deprecated. Use packaging.version instead.

```
return LooseVersion(v) >= LooseVersion(check)
```

Failed to detect the name of this notebook, you can set it manually with the
WANDB_NOTEBOOK_NAME environment variable to enable code saving.

Create sweep with ID: ofvxspxo
Sweep URL: [https://wandb.ai/me19b190/Hyper%20parameter%20tuning%20RL%20full%20re
turn%20-%20v1/sweeps/ofvxspxo](https://wandb.ai/me19b190/Hyper%20parameter%20tuning%20RL%20full%20re%20turn%20-%20v1/sweeps/ofvxspxo)

```
[ ]: wandb.agent(sweep_id, train, count=3)
```

wandb: Agent Starting Run: 0g4dbjhd with config:

wandb: LR: 0.0004888636786360583

wandb: NUM_EPISODES: 2000

wandb: NUM_LAYERS: 8

wandb: NUM_NEURONS_EACH_LAYER: 128

Failed to detect the name of this notebook, you can set it manually with the
WANDB_NOTEBOOK_NAME environment variable to enable code saving.

wandb: Currently logged in as: **me19b190**. Use `wandb

login --relogin` to force relogin

C:\Users\CFI Workstation\AppData\Local\Packages\PythonSoftwareFoundation.Python.
3.10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\wandb\sdk\lib\ipython.py:47: DeprecationWarning: Importing display from
IPython.core.display is deprecated since IPython 7.14, please import from
IPython display

from IPython.core.display import HTML, display # type: ignore

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

C:\Users\CFI Workstation\AppData\Local\Packages\PythonSoftwareFoundation.Python.
3.10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\gym\core.py:317: DeprecationWarning: **WARN: Initializing wrapper in
old step API which returns one bool instead of two. It is recommended to set
`new_step_api=True` to use new step API. This will be the default behaviour in
future.**

deprecation(
deprecation(
C:\Users\CFI Workstation\AppData\Local\Packages\PythonSoftwareFoundation.Python.
3.10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\gym\wrappers\step_api_compatibility.py:39: DeprecationWarning:

**WARN: Initializing environment in old step API which returns one bool
instead of two. It is recommended to set `new_step_api=True` to use new step
API. This will be the default behaviour in future.**

deprecation(
deprecation(
C:\Users\CFI Workstation\AppData\Local\Packages\PythonSoftwareFoundation.Python.
3.10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\gym\wrappers\step_api_compatibility.py:39: DeprecationWarning:

```
C:\Users\CFI Workstation\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\gym\core.py:256: DeprecationWarning: WARN: Function
```

```
`env.seed(seed)` is marked as deprecated and will be removed in the future.
```

```
Please use `env.reset(seed=seed)` instead.
```

```
deprecation(
```

```
Run:  1  -----
Episode  10 Reward -500.000000 Average Reward -500.000000
Episode  20 Reward -500.000000 Average Reward -500.000000
Episode  30 Reward -500.000000 Average Reward -500.000000
Episode  40 Reward -500.000000 Average Reward -500.000000
Episode  50 Reward -500.000000 Average Reward -500.000000
```

```
[ ]: #execution interrupted due to bad hyperparameter setting
```

```
steps_over_10_runs = []
rewards_over_10_runs = []
episodes = []
min_len = np.inf
for run in range(1):
    print("Run: ",run+1,"_")
    ↪-----")
    env = gym.make('CartPole-v1')

    seed = np.random.randint(low=0, high = 100)
    env.seed(seed)
    state_shape = env.observation_space.shape[0]
    action_shape = env.action_space.n

    begin_time = datetime.datetime.now()     #(num_layers, ↪
    ↪num_neurons_each_layer, LR,num_episodes,seed):
    rewards_each_ep, steps_each_ep, episodes_solved = ↪
    ↪Actor_critic_one_step_return(NUM_LAYERS, NUM_NEURONS_EACH_LAYER, ↪
    ↪LR,NUM_EPISODES,seed)
    steps_over_10_runs.append(steps_each_ep)
    rewards_over_10_runs.append(rewards_each_ep)
    episodes.append(episodes_solved)

    if len(rewards_each_ep) < min_len:
        min_len = len(rewards_each_ep)
    time_taken = datetime.datetime.now() - begin_time

    print(time_taken)

rewards = np.array(rewards_over_10_runs)
```

```

avg_episodes = int(np.array(episodes).mean())
steps = np.array(steps_over_10_runs)

y, error_reward = calculate_mean(rewards[:avg_episodes])
plt.figure(figsize = (15,5))
plt.subplot(121)
    # for i in range(len(rewards_over_10_runs)):
    #     plt.plot(np.
    #         range(len(rewards_over_10_runs[i])),rewards_over_10_runs[i],label='run_
    #         '+str(i))
plt.plot(np.arange(len(y))+1, y, color='darkblue',label='average')
plt.title("Episodes vs rewards")
plt.xlabel("Number of episodes")
plt.ylabel("Reward per episode")
plt.legend()
plt.subplot(122)
y, error = calculate_mean(steps[:avg_episodes])
plt.plot(np.arange(len(y))+1, y, color='darkblue',label='average')
plt.title("Episodes vs steps ")
plt.xlabel("Number of episodes")
plt.ylabel("Steps taken per episode")
plt.legend()
plt.subplot(133)
plt.plot(np.arange(len(error_reward))+1, error_reward , color='darkblue',label=
    'Deviation of rewards') #plotting the variance = std squared
plt.title("Episodes vs Variance ")
plt.xlabel("Number of episodes")
plt.ylabel("Observed Variance")
plt.legend()
plt.savefig('hyp_cartpole_ac_one_step_return'+str(NUM_EPISODES)+str(LR)+'.jpg',
    dpi = 250)
plt.show()

for i in range(avg_episodes):
    print("Episode: ",i+1, 'Number of steps: ',round(y[i]))

```

```

Run: 1 -----
Episode 10 Reward -205.000000 Average Reward -252.300000
Episode 20 Reward -141.000000 Average Reward -193.500000
Episode 30 Reward -220.000000 Average Reward -173.200000
Episode 40 Reward -163.000000 Average Reward -167.200000
Episode 50 Reward -229.000000 Average Reward -150.100000
Episode 60 Reward -169.000000 Average Reward -158.100000
Episode 70 Reward -140.000000 Average Reward -158.200000
Episode 80 Reward -232.000000 Average Reward -133.200000
Episode 90 Reward -500.000000 Average Reward -316.100000

```

```

Episode 100 Reward -500.000000 Average Reward -500.000000
Episode 110 Reward -500.000000 Average Reward -500.000000
Episode 120 Reward -500.000000 Average Reward -500.000000
Episode 130 Reward -500.000000 Average Reward -500.000000
Episode 140 Reward -500.000000 Average Reward -500.000000
Episode 150 Reward -500.000000 Average Reward -500.000000
Episode 160 Reward -500.000000 Average Reward -500.000000
Episode 170 Reward -500.000000 Average Reward -500.000000
Episode 180 Reward -500.000000 Average Reward -500.000000
Episode 190 Reward -500.000000 Average Reward -500.000000
Episode 200 Reward -500.000000 Average Reward -500.000000
Episode 210 Reward -500.000000 Average Reward -500.000000
Episode 220 Reward -500.000000 Average Reward -500.000000
Episode 230 Reward -500.000000 Average Reward -500.000000
Episode 240 Reward -500.000000 Average Reward -500.000000
Episode 250 Reward -500.000000 Average Reward -500.000000
Episode 260 Reward -500.000000 Average Reward -500.000000

```

KeyboardInterrupt Traceback (most recent call last)

Cell In[11], line 15

```

    12 action_shape = env.action_space.n
    14 begin_time = datetime.datetime.now()    #(num_layers,
↳ num_neurons_each_layer, LR, num_episodes, seed):
--> 15 rewards_each_ep, steps_each_ep, episodes_solved =
↳ Actor_critic_one_step_return(NUM_LAYERS, NUM_NEURONS_EACH_LAYER,
↳ LR, NUM_EPISODES, seed)
    16 steps_over_10_runs.append(steps_each_ep)
    17 rewards_over_10_runs.append(rewards_each_ep)

```

Cell In[6], line 31, in Actor_critic_one_step_return(num_layers,
↳ num_neurons_each_layer, LR, num_episodes, seed)

```

    29 next_state = next_state.reshape(1,-1)
    30 ep_rew += reward    ##Updating episode reward
--> 31 agent.learn(state, action, reward, next_state, done) ##Update
↳ Parameters
    32 state = next_state ##Updating State
    33 reward_list.append(ep_rew)

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

↳ 10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\tensorflow\python\util\
↳ py:150, in filter_traceback.<locals>.error_handler(*args, **kwargs)
    148 filtered_tb = None
    149 try:
--> 150 return fn(*args, **kwargs)
    151 except Exception as e:
    152 filtered_tb = _process_traceback_frames(e.__traceback__)

```

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\tensorflow\python\eager\
↳py:894, in Function.__call__(self, *args, **kwargs)
    891 compiler = "xla" if self._jit_compile else "nonXla"
    893 with OptionalXlaContext(self._jit_compile):
--> 894     result = self._call(*args, **kwargs)
    896 new_tracing_count = self.experimental_get_tracing_count()
    897 without_tracing = (tracing_count == new_tracing_count)

```

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\tensorflow\python\eager\
↳py:926, in Function._call(self, *args, **kwargs)
    923 self._lock.release()
    924 # In this case we have created variables on the first call, so we run
↳the
    925 # defunned version which is guaranteed to never create variables.
--> 926 return self._no_variable_creation_fn(*args, **kwargs) # pylint:
↳disable=not-callable
    927 elif self._variable_creation_fn is not None:
    928     # Release the lock early so that multiple threads can perform the cal
    929     # in parallel.
    930     self._lock.release()

```

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\tensorflow\python\eager\
↳py:143, in TracingCompiler.__call__(self, *args, **kwargs)
    140 with self._lock:
    141     (concrete_function,
    142      filtered_flat_args) = self._maybe_define_function(args, kwargs)
--> 143 return concrete_function._call_flat(
    144     filtered_flat_args, captured_inputs=concrete_function.
↳captured_inputs)

```

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\tensorflow\python\eager\
↳py:1757, in ConcreteFunction._call_flat(self, args, captured_inputs,
↳cancellation_manager)
    1753 possible_gradient_type = gradients_util.PossibleTapeGradientTypes(args)
    1754 if (possible_gradient_type == gradients_util.POSSIBLE_GRADIENT_TYPES_NO_E
    1755     and executing_eagerly):
    1756     # No tape is watching; skip to running the function.
-> 1757 return self._build_call_outputs(self._inference_function.call(
    1758     ctx, args, cancellation_manager=cancellation_manager))
    1759 forward_backward = self._select_forward_and_backward_functions(
    1760     args,
    1761     possible_gradient_type,
    1762     executing_eagerly)
    1763 forward_function, args_with_tangents = forward_backward.forward()

```

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
  ↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\tensorflow.python\eager
  ↳py:381, in _EagerDefinedFunction.call(self, ctx, args, cancellation_manager)
    379 with _InterpolateFunctionError(self):
    380     if cancellation_manager is None:
--> 381         outputs = execute.execute(
    382             str(self.signature.name),
    383             num_outputs=self._num_outputs,
    384             inputs=args,
    385             attrs=attrs,
    386             ctx=ctx)
    387     else:
    388         outputs = execute.execute_with_cancellation(
    389             str(self.signature.name),
    390             num_outputs=self._num_outputs,
    (...)
    393             ctx=ctx,
    394             cancellation_manager=cancellation_manager)

```

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
  ↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\tensorflow.python\eager
  ↳py:52, in quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
    50 try:
    51     ctx.ensure_initialized()
---> 52     tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name
    53                                         inputs, attrs, num_outputs)
    54 except core._NotOkStatusException as e:
    55     if name is not None:

```

KeyboardInterrupt: