

pa2-dqn-mountain-car

March 29, 2023

```
[ ]: import math
import numpy as np
import random
import torch
import torch.nn as nn
import torch.nn.functional as F
from collections import namedtuple, deque
import torch.optim as optim
import datetime
import gym
from gym.wrappers.record_video import RecordVideo
import glob
import io
import base64
import matplotlib.pyplot as plt
from IPython.display import HTML
from pyvirtualdisplay import Display
import tensorflow as tf
from IPython import display as ipythondisplay
from PIL import Image
import tensorflow_probability as tfp
import warnings
warnings.filterwarnings("ignore")
```

```
[ ]: env = gym.make('MountainCar-v0')
env.seed(0)

state_shape = env.observation_space.shape[0]
no_of_actions = env.action_space.n

print(state_shape)
print(no_of_actions)
print(env.action_space.sample())
print("-----")

state = env.reset()
```

```

print(state)
print("----")

action = env.action_space.sample()

print(action)
print("----")

next_state, reward, done, info = env.step(action)

print(next_state)
print(reward)
print(done)
print(info)
print("----")

```

```

2
3
1
----
[-0.47260767  0.          ]
----
1
----
[-4.7298861e-01 -3.8094356e-04]
-1.0
False
{}
----

```

```

[ ]: sweep_config = {
    'method': 'bayes'
}

```

```

[ ]: metric = {
    'name' : 'num_episodes_tosolve_cartpole',
    'goal' : 'minimize'
}
sweep_config['metric'] = metric

```

```

[ ]: parameters_dict = {
    'NUM_NEURONS_EACH_LAYER' : {
        'values': [64,32,16]
    },
    'NUM_LAYERS' :{
        'values': [2,4,8]
    }
}

```

```

    },
  }
  sweep_config['parameters'] = parameters_dict

```

```

[ ]: import math

parameters_dict.update({
    'LR': {

        'distribution': 'uniform',
        'min': (0.0001),
        'max': (0.01)
    },
    'BATCH_SIZE': {

        'distribution': 'q_log_uniform_values',
        'q' : 1,
        'min': (64),
        'max': (512),
    },
    'UPDATE_EVERY': {

        'distribution': 'q_log_uniform_values',
        'q' : 1,
        'min': (10),
        'max': (200),
    },
    'GAMMA': {

        'distribution': 'uniform',
        'min': 0.8,
        'max': 1,
    },
    'MAX_TRUNCATION': {

        'distribution': 'uniform',
        'min': (0.5),
        'max': (2),
    },
    'EPS_DECAY': {

        'distribution': 'uniform',
        'min': 0.9,
        'max': 0.995,
    },
    'a_TUNE': {

```

```

        'distribution': 'uniform',
        'min': 0,
        'max': 1,
    }

})

```

```
[ ]: import pprint
```

```
pprint.pprint(sweep_config)
```

```

{'method': 'bayes',
 'metric': {'goal': 'minimize', 'name': 'num_episodes_tosolve_cartpole'},
 'parameters': {'BATCH_SIZE': {'distribution': 'q_log_uniform_values',
                                'max': 512,
                                'min': 64,
                                'q': 1},
                 'EPS_DECAY': {'distribution': 'uniform',
                                'max': 0.995,
                                'min': 0.9},
                 'GAMMA': {'distribution': 'uniform', 'max': 1, 'min': 0.8},
                 'LR': {'distribution': 'uniform', 'max': 0.01, 'min': 0.0001},
                 'MAX_TRUNCATION': {'distribution': 'uniform',
                                    'max': 2,
                                    'min': 0.5},
                 'NUM_LAYERS': {'values': [2, 4, 8]},
                 'NUM_NEURONS_EACH_LAYER': {'values': [64, 32, 16]},
                 'UPDATE_EVERY': {'distribution': 'q_log_uniform_values',
                                   'max': 200,
                                   'min': 10,
                                   'q': 1},
                 'a_TUNE': {'distribution': 'uniform', 'max': 1, 'min': 0}}}

```

```
[ ]: sweep_id = wandb.sweep(sweep_config, project="Hyper parameter tuning Cartpole -v1")
```

Failed to detect the name of this notebook, you can set it manually with the WANDB_NOTEBOOK_NAME environment variable to enable code saving.

Create sweep with ID: 96xha17k

Sweep URL: <https://wandb.ai/me19b190/Hyper%20parameter%20tuning%20Cartpole%20-%20v1/sweeps/96xha17k>

```
[ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
```

```

class QNetwork1(nn.Module): #take state s and outputs Q(s,a)

    def __init__(self,seed, state_size, action_size,
↳num_layers,neurons_each_layer): #num_layers = number of hidden
↳layers,neurons_each_layer = number of neurons in hidden layers
        super(QNetwork1,self).__init__()
        self.seed = torch.manual_seed(seed)
        self.input_size = state_size
        self.output_size = action_size
        self.num_layers = num_layers
        self.neurons_each_layer = neurons_each_layer
        self.linears = nn.ModuleList([nn.Linear(self.input_size, self.
↳neurons_each_layer[0])])
        self.linears.extend([nn.Linear(self.neurons_each_layer[i-1], self.
↳neurons_each_layer[i]) for i in range(1, self.num_layers)])
        self.linears.append(nn.Linear(self.neurons_each_layer[-1], self.
↳output_size))

    def forward(self,state):
        h=None
        for i in range(self.num_layers+1):
            if i == 0:
                h = F.relu(self.linears[0](state))
            elif i<(self.num_layers):
                h = F.relu(self.linears[i](h))
            else:
                return self.linears[self.num_layers](h)

```

```

[ ]: import random
import torch
import numpy as np
from collections import deque, namedtuple

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

class ReplayBuffer:
    """Fixed-size buffer to store experience tuples."""

    def __init__(self, action_size, buffer_size, batch_size, seed):
        """Initialize a ReplayBuffer object.

        Params
        =====
        action_size (int): dimension of each action
        buffer_size (int): maximum size of buffer
        batch_size (int): size of each training batch
        seed (int): random seed

```

```

        """
        self.action_size = action_size
        self.memory = deque(maxlen=buffer_size)
        self.batch_size = batch_size
        self.experience = namedtuple("Experience", field_names=["state",
↪ "action", "reward", "next_state", "done"])
        self.seed = random.seed(seed)

    def add(self, state, action, reward, next_state, done): #added TD
        """Add a new experience to memory."""
        e = self.experience(state, action, reward, next_state, done)
        self.memory.append(e)

    def sample(self, batch_size=None, a=1.0):
        """Randomly sample a batch of experiences from memory."""
        experiences = random.sample(self.memory, k=self.batch_size)

        states = torch.from_numpy(np.vstack([e.state for e in experiences if e
↪ is not None])).float().to(device)
        actions = torch.from_numpy(np.vstack([e.action for e in experiences if
↪ e is not None])).long().to(device)
        rewards = torch.from_numpy(np.vstack([e.reward for e in experiences if
↪ e is not None])).float().to(device)
        next_states = torch.from_numpy(np.vstack([e.next_state for e in
↪ experiences if e is not None])).float().to(device)
        dones = torch.from_numpy(np.vstack([e.done for e in experiences if e is
↪ not None])).astype(np.uint8).float().to(device)

        return (states, actions, rewards, next_states, dones)

    def __len__(self):
        """Return the current size of internal memory."""
        return len(self.memory)

```

```

[ ]: def calculate_mean(array):
    lens = [len(i) for i in array]
    arr = np.ma.empty((np.max(lens), len(array)))
    arr.mask = True
    for idx, l in enumerate(array):
        arr[:len(l), idx] = l
    return arr.mean(axis = -1), arr.std(axis=-1)

```

```

[ ]: ''' Defining DQN Algorithm '''
state_shape = env.observation_space.shape[0]
action_shape = env.action_space.n

```

```

def dqn(agent,n_episodes=5000, max_t=500, eps_start=1.0, eps_end=0.01,
↪eps_decay=0.999,a = 0.7):

    scores = []
    ''' list containing scores from each episode '''
    steps = []
    scores_window_printing = deque(maxlen=10)
    ''' For printing in the graph '''
    scores_window= deque(maxlen=100)
    ''' last 100 scores for checking if the avg is more than 195 '''
    scores_exit_bad_paramas = deque(maxlen=200)
    eps = eps_start
    ''' initialize epsilon '''

    for i_episode in range(1, n_episodes+1):
        state = env.reset()
        score = 0
        step = 0
        for t in range(max_t):
            state_adj = (state-env.observation_space.low)*np.array([10,50])
            state_adj = np.round(state_adj,0).astype(int)
            action = agent.act(state_adj, eps)
            next_state, reward, done, _ = env.step(action)
            next_state_adj = (next_state-env.observation_space.low)*np.
↪array([10,50])
            next_state_adj = np.round(next_state_adj,0).astype(int)
            agent.step(state_adj, action, reward, next_state_adj, done,a = a)
            state = next_state
            score += reward
            step += 1
            if done:
                break
        steps.append(step)
        scores_window.append(score)
        scores_window_printing.append(score)
        scores_exit_bad_paramas.append(score)

        ''' save most recent score '''

        eps = max(eps_end, eps_decay*eps)
        ''' decrease epsilon '''

        print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.
↪mean(scores_window)), end="")
        if i_episode % 10 == 0:
            scores.append(np.mean(scores_window_printing))

```

```

        if i_episode % 100 == 0:
            print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.
↪mean(scores_window)))
            if np.mean(scores_window)>=-110.0:
                print('\nEnvironment solved in {:d} episodes! \tAverage Score: {:.
↪2f}'.format(i_episode-100, np.mean(scores_window)))
                break
            return [np.array(steps), np.array(scores), i_episode, False]

''' Trial run to check if algorithm runs and saves the data '''
BATCH_SIZE= 64
EPS_DECAY= 0.995
GAMMA= 0.99
LR= 5e-4
MAX_TRUNCATION= 1
NUM_LAYERS= 2
NUM_NEURONS_EACH_LAYER= 64
UPDATE_EVERY= 20
a_TUNE= 0

agent_l = None
steps_over_10_runs = []
rewards_over_10_runs = []
episodes = []
min_len = np.inf
for run in range(1):
    print("Run: ", run+1, "\n
↪-----")
    env = gym.make('MountainCar-v0', max_episode_steps=500)
    seed = run
    env.seed(seed)
    state_shape = env.observation_space.shape[0]
    action_shape = env.action_space.n

    begin_time = datetime.datetime.now()
    neurons_each_layer_ = [NUM_NEURONS_EACH_LAYER]*NUM_LAYERS
    agent = TutorialAgent(state_size= state_shape, action_size=action_shape,
↪seed=seed, num_layers=NUM_LAYERS, neurons_each_layer=neurons_each_layer_
↪, lr=LR, gamma=GAMMA, update_every=UPDATE_EVERY, batch_size=BATCH_SIZE, max_tRUNCATION=MAX_TRUNC
↪steps_eps_greedy, reward_eps_greedy, episodes_eps_greedy, agent_l =
↪dqn(agent, eps_decay = EPS_DECAY, a = a_TUNE)
    steps_over_10_runs.append(steps_eps_greedy)
    rewards_over_10_runs.append(reward_eps_greedy)
    episodes.append(episodes_eps_greedy)
    if len(reward_eps_greedy) < min_len:
        min_len = len(reward_eps_greedy)
    time_taken = datetime.datetime.now() - begin_time

```



```

print(time_taken)

rewards = np.array(rewards_over_10_runs)
avg_episodes = int(np.array(episodes).mean())
steps = np.array(steps_over_10_runs)

```

```

Run:  1  -----
Episode 100      Average Score: -500.00
Episode 200      Average Score: -500.00
Episode 300      Average Score: -492.10
Episode 400      Average Score: -403.26
Episode 500      Average Score: -251.92
Episode 600      Average Score: -169.75
Episode 700      Average Score: -147.15
Episode 800      Average Score: -142.20
Episode 900      Average Score: -142.17
Episode 1000     Average Score: -150.11
Episode 1100     Average Score: -138.30
Episode 1200     Average Score: -136.35
Episode 1300     Average Score: -133.42
Episode 1400     Average Score: -136.04
Episode 1500     Average Score: -134.69
Episode 1600     Average Score: -140.36
Episode 1700     Average Score: -140.18
Episode 1800     Average Score: -147.31
Episode 1900     Average Score: -147.55
Episode 2000     Average Score: -140.56
Episode 2100     Average Score: -141.59
Episode 2200     Average Score: -135.20
Episode 2300     Average Score: -144.09
Episode 2400     Average Score: -134.94
Episode 2500     Average Score: -132.15
Episode 2600     Average Score: -142.41
Episode 2700     Average Score: -146.01
Episode 2800     Average Score: -149.58
Episode 2900     Average Score: -148.96
Episode 3000     Average Score: -149.31
Episode 3100     Average Score: -149.58
Episode 3200     Average Score: -148.10
Episode 3300     Average Score: -154.22
Episode 3400     Average Score: -153.82
Episode 3500     Average Score: -154.79
Episode 3600     Average Score: -143.06
Episode 3700     Average Score: -139.03
Episode 3800     Average Score: -143.84
Episode 3900     Average Score: -153.15

```

```

Episode 4000    Average Score: -146.50
Episode 4100    Average Score: -153.11
Episode 4200    Average Score: -169.81
Episode 4300    Average Score: -159.53
Episode 4400    Average Score: -154.84
Episode 4500    Average Score: -157.75
Episode 4600    Average Score: -154.27
Episode 4700    Average Score: -152.30
Episode 4800    Average Score: -142.60
Episode 4900    Average Score: -143.62
Episode 5000    Average Score: -142.63
1:08:07.755923

```

Hyperparameters: 1. Number of neurons in neural network 64,128,256,512 2. Number of layers in neural network 2,4,8,10 3. learning rate: `log_uniform('1e-6' to '1e-1')` 4. `buffer_Size = log_uniform(10-300)` 5. `update_frequency = log_uniform(20,200)` 6. `truncation_Limit = log_uniform(100,1000,10000)` 7. discount factor = `uniform(0,1)` 8. epsilon decay = `uniform(0.9,0.995)` 9. `a = uniform(0.01,1)` 10. `b = similar to epsilon decay of epsilon` Try 12 different configs

```
[ ]: rewards
```

```

[ ]: array([[ -500. , -500. , -500. , -500. , -500. , -500. , -500. , -500. ,
           -500. , -500. , -500. , -500. , -500. , -500. , -500. , -500. ,
           -500. , -500. , -500. , -500. , -497.8, -500. , -500. , -467.1,
           -496.4, -500. , -500. , -488.8, -471.5, -499.4, -489.6, -402.4,
           -436.6, -425.8, -379.5, -454.3, -404.2, -384.3, -339.4, -316.5,
           -315.8, -274.6, -290.7, -257.5, -268.7, -273.7, -206. , -218. ,
           -213.4, -200.8, -179.7, -175.2, -160.4, -177.1, -166.7, -169.4,
           -155.4, -183.2, -171.8, -158.6, -155.9, -168.8, -149.3, -133.5,
           -155.8, -146.9, -129.5, -144.1, -140. , -147.7, -142.7, -137.8,
           -150. , -144.1, -124.7, -152.4, -153.4, -141.3, -149. , -126.6,
           -142. , -146.8, -137.1, -141.4, -124.1, -142.7, -153. , -143. ,
           -142.2, -149.4, -140. , -144.6, -147.4, -154.1, -166.1, -140.1,
           -150.2, -148.4, -154.4, -155.8, -135. , -141. , -147.2, -151.6,
           -143. , -119.1, -137.7, -129.4, -139. , -140. , -145.7, -139.3,
           -122.6, -138.5, -137.9, -150.4, -138.5, -131. , -133.2, -126.4,
           -126.4, -126.3, -147.4, -139.5, -130.6, -128.5, -135.1, -137.6,
           -135. , -127.8, -130.3, -133.1, -123.2, -155.2, -136.4, -141.6,
           -132. , -114.5, -149.7, -144.4, -120.8, -124.2, -125.7, -128.3,
           -150.3, -146.1, -160.7, -129.4, -136.2, -125.2, -146.1, -151.1,
           -143.1, -139.2, -131. , -139.8, -140.1, -149.1, -141.2, -122.9,
           -128.1, -169.3, -131.7, -155. , -131.4, -140.4, -123.9, -133.9,
           -142.8, -145.3, -167. , -158.7, -142.2, -152.2, -138.6, -146. ,
           -143.1, -146. , -129.7, -149.6, -123.7, -153.3, -164.3, -143.5,
           -146.6, -156.6, -139. , -144. , -147.9, -156.6, -128.1, -140.2,
           -141.9, -121.1, -158.1, -151.9, -127.4, -140.8, -149.3, -146.8,
           -138.8, -137.2, -150.2, -141.1, -140.3, -149. , -142.1, -150.8,

```

```

-136.8, -129.6, -144.6, -134.1, -153.3, -140.2, -119.4, -130.6,
-137.2, -130.7, -121. , -140.9, -137. , -148.7, -165.7, -152.1,
-139.7, -136.4, -145.1, -136.8, -137.3, -142.1, -127.7, -128.4,
-145.1, -136.1, -121.7, -136.2, -145.4, -128.3, -134.5, -146. ,
-130.9, -132.2, -128.4, -151.6, -126.9, -127.6, -130.8, -126.3,
-126.2, -140.6, -124. , -135. , -136.2, -145.4, -127.6, -153. ,
-145.6, -151.2, -162.5, -143.6, -143.3, -140. , -153.7, -126.9,
-145.9, -154.4, -156.3, -165.9, -137.9, -135.8, -163.2, -157.7,
-149.2, -153.4, -137.3, -146.8, -152.4, -129.7, -137.5, -168.6,
-153.5, -153.9, -160.3, -153.1, -140.2, -154.2, -138.7, -150.6,
-139. , -146.1, -152.7, -149.9, -145. , -149.3, -156.5, -156.7,
-161.7, -143.2, -124.5, -153.6, -152.5, -142.4, -148. , -144.7,
-148.4, -158.5, -161.7, -146.6, -144.8, -148.2, -155.6, -150.1,
-151.5, -137.6, -144.7, -151.9, -140.9, -144.6, -149.7, -154.4,
-147.8, -153.1, -152.7, -144.2, -151.5, -146. , -168.4, -159.4,
-167.1, -152. , -174.6, -143.8, -165.5, -159.7, -151.3, -148. ,
-140.8, -147.9, -144.8, -161.8, -136.2, -174.1, -165.3, -167.4,
-143.3, -156.4, -143. , -167.4, -155.6, -139.2, -148.7, -142.4,
-140. , -136.3, -146.9, -153.8, -146.1, -143.8, -129. , -143.6,
-147.4, -143.8, -152.9, -124.9, -131.3, -132.5, -140. , -142.8,
-156.3, -118.4, -140.9, -151.9, -138.7, -144.5, -147.6, -138.4,
-143.7, -146. , -137.6, -149.1, -165.9, -131.5, -158.8, -169.4,
-150.7, -158.4, -155.6, -130.1, -155.3, -155.8, -135. , -159.9,
-143.7, -147.5, -132.9, -144.6, -161.6, -157.1, -140.4, -142.3,
-154.9, -142.1, -160. , -151.3, -140.5, -150.3, -165.1, -165.2,
-148.9, -152.8, -153.6, -172.9, -160.7, -173.2, -164.7, -160.6,
-167.5, -163.5, -212.4, -169. , -179.2, -150. , -166.5, -165.7,
-144. , -153.9, -165.3, -176.5, -150.1, -144.1, -157.7, -138.9,
-150.9, -174.3, -149. , -148.8, -162.9, -160.1, -142.2, -163.6,
-146.1, -152.3, -175.4, -162.1, -154.8, -143.3, -180.3, -168.6,
-150.8, -143.8, -156.6, -151.1, -160.6, -154.8, -162.1, -161.7,
-150.2, -157. , -141.1, -147.5, -149.7, -174. , -156.2, -154.4,
-149.4, -125.9, -154.2, -157.6, -142.7, -158.9, -148. , -140.6,
-158.7, -125.8, -143.7, -145.2, -143.3, -143.2, -134.9, -142.6,
-135.5, -128.5, -143.7, -139.3, -152. , -146.1, -136.6, -144.8,
-161.2, -148.5, -129.4, -134.5, -136.4, -151.6, -142.6, -146.3,
-152. , -143.2, -151.7, -138.6]])

```

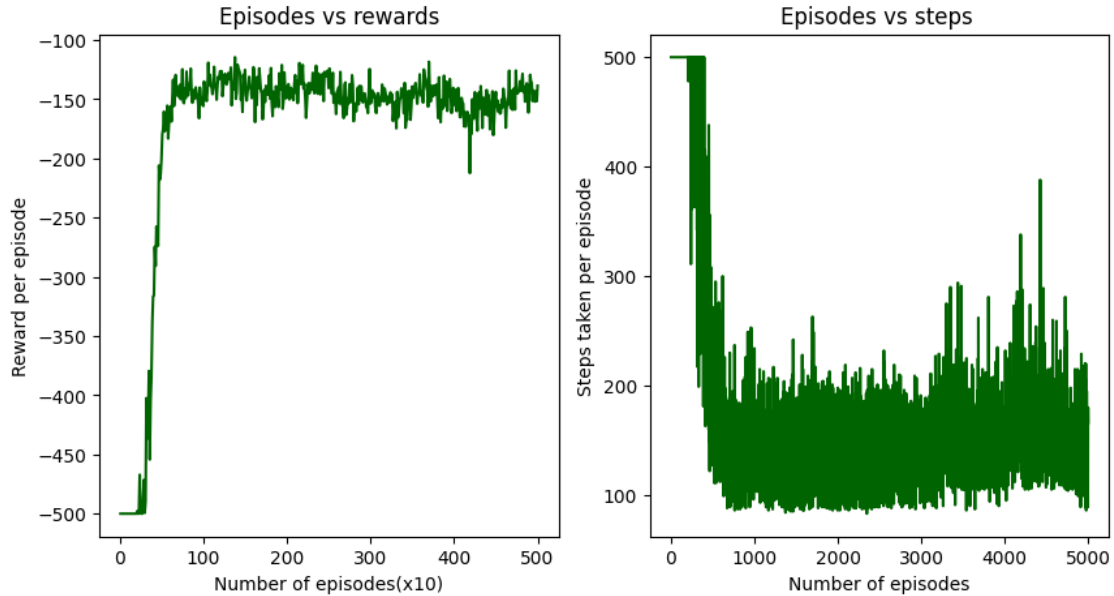
```

[ ]: plt.figure(figsize = (10,5))
plt.subplot(121)
plt.plot(np.arange(len(rewards[0]))+1, rewards[0], color='darkgreen')
plt.title("Episodes vs rewards")
plt.xlabel("Number of episodes(x10)")
plt.ylabel("Reward per episode")

plt.subplot(122)
plt.plot(np.arange(len(steps[0]))+1, steps[0], color='darkgreen')

```

```
plt.title("Episodes vs steps ")
plt.xlabel("Number of episodes")
plt.ylabel("Steps taken per episode")
plt.show()
```



```
[ ]: pprint.pprint(sweep_config)
```

```
{'method': 'random',
 'metric': {'goal': 'minimize', 'name': 'num_episodes_tosolve_cartpole'},
 'parameters': {'BATCH_SIZE': {'distribution': 'q_log_uniform_values',
                                'max': 512,
                                'min': 64,
                                'q': 1},
                 'EPS_DECAY': {'distribution': 'uniform',
                                'max': 0.995,
                                'min': 0.9},
                 'GAMMA': {'distribution': 'uniform', 'max': 1, 'min': 0.8},
                 'LR': {'distribution': 'uniform', 'max': 0.01, 'min': 0.0001},
                 'MAX_TRUNCATION': {'distribution': 'uniform',
                                    'max': 2,
                                    'min': 0.5},
                 'NUM_LAYERS': {'values': [2, 4, 8]},
                 'NUM_NEURONS_EACH_LAYER': {'values': [64, 32, 16]},
                 'UPDATE_EVERY': {'distribution': 'q_log_uniform_values',
                                   'max': 200,
                                   'min': 10,
                                   'q': 1},
```

```
'a_TUNE': {'distribution': 'uniform', 'max': 1, 'min': 0}}
```

```
[ ]: def train(config = None):
    with wandb.init(config = config):
        config = wandb.config
        BUFFER_SIZE = int(1e5) # replay buffer size
        BATCH_SIZE = config.BATCH_SIZE ## minibatch size
        GAMMA = config.GAMMA ## discount factor
        LR = config.LR ## learning rate
        UPDATE_EVERY = config.UPDATE_EVERY ## how often to update the network
        ↪(When Q target is present)
        NUM_NEURONS_EACH_LAYER = config.NUM_NEURONS_EACH_LAYER ##number of
        ↪neurons in each hidden layer
        NUM_LAYERS = config.NUM_LAYERS ##number of layers for the
        ↪neural network
        MAX_TRUNCATION = config.MAX_TRUNCATION ##max number of
        ↪steps in each episode(needed only for cartpole, acrobot and mountain car
        ↪these are already defined)
        EPS_DECAY = config.EPS_DECAY ##epsilon decay for
        ↪epsilon exploration
        a_TUNE = config.a_TUNE

        steps_over_10_runs = []
        rewards_over_10_runs = []
        episodes = []
        min_len = np.inf
        bad = False
        for run in range(10):
            print("Run: ",run,"
            ↪-----")
            env = gym.make('MountainCar-v0')

            seed = run
            env.seed(seed)
            state_shape = env.observation_space.shape[0]
            action_shape = env.action_space.n

            begin_time = datetime.datetime.now()
            agent = TutorialAgent(state_size=state_shape,action_size =
            ↪action_shape,seed =
            ↪seed,num_layers=NUM_LAYERS,neurons_each_layer=[NUM_NEURONS_EACH_LAYER]*NUM_LAYERS,lr=
            ↪LR,gamma= GAMMA,update_every=UPDATE_EVERY,batch_size =
            ↪BATCH_SIZE,max_truncation = MAX_TRUNCATION)
            steps_eps_greedy,reward_eps_greedy, episodes_eps_greedy,bad =
            ↪dqn(agent,eps_decay = EPS_DECAY, a = a_TUNE)
            steps_over_10_runs.append(steps_eps_greedy)
            rewards_over_10_runs.append(reward_eps_greedy)
```

```

    episodes.append(episodes_eps_greedy)
    if len(reward_eps_greedy) < min_len:
        min_len = len(reward_eps_greedy)
    time_taken = datetime.datetime.now() - begin_time

    print(time_taken)

rewards = np.array(rewards_over_10_runs)
avg_episodes = int(np.array(episodes).mean())
steps = np.array(steps_over_10_runs)

y, error = calculate_mean(rewards)
plt.figure(figsize = (15,10))
plt.subplot(121)
plt.axvline(x = avg_episodes/10 , color = 'black', label = 'Environment_
↳solved')
plt.scatter(avg_episodes/10,0, marker = 'x')
for i in range(len(rewards_over_10_runs)):
    plt.plot(np.
↳arange(len(rewards_over_10_runs[i])),rewards_over_10_runs[i],label='run_
↳'+str(i))
    plt.plot(np.arange(len(y))+1, y, color='darkblue',label='average')
    plt.fill_between(np.arange(len(y))+1, y-error, y+error,color =_
↳'lightskyblue',label = 'standard deviation')
plt.title("Episodes vs rewards")
plt.xlabel("Number of episodes(x10)")
plt.ylabel("Reward per episode")
plt.legend()
plt.subplot(122)
y, error = calculate_mean(steps)
for i in range(len(steps_over_10_runs)):
    plt.plot(np.
↳arange(len(steps_over_10_runs[i])),steps_over_10_runs[i],label='run '+str(i))
    plt.plot(np.arange(len(y))+1, y, color='darkblue',label='average')
    plt.fill_between(np.arange(len(y))+1, y-error, y+error,color =_
↳'lightskyblue',label = 'standard deviation')
plt.title("Episodes vs steps ")
plt.xlabel("Number of episodes")
plt.ylabel("Steps taken per episode")
plt.legend()
plt.savefig(str(CONFIGGGG)+'runhyp'+str(LR)+'.jpg', dpi = 250)
plt.show()
for i in range(avg_episodes):
    print("Episode: ",i+1, 'Number of steps: ',round(y[i]))
if bad == False:
    wandb.log({"num_episodes_tosolve_cartpole":avg_episodes-100})
else:

```

```
wandb.log({"num_episodes_tosolve_cartpole":5000})
```

```
[ ]: wandb.agent(sweep_id, train, count=12)
```

wandb: Agent Starting Run: btvtcodk with config:

wandb: BATCH_SIZE: 319

wandb: EPS_DECAY: 0.902341235943596

wandb: GAMMA: 0.8144181085986043

wandb: LR: 0.0022755859727504983

wandb: MAX_TRUNCATION: 1.074217920563914

wandb: NUM_LAYERS: 2

wandb: NUM_NEURONS_EACH_LAYER: 64

wandb: UPDATE_EVERY: 15

wandb: a_TUNE: 0.862186399649838

wandb: Currently logged in as: **me19b190**. Use ``wandb`

`login --relogin`` to force relogin

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Run: 0 -----

Episode 100 Average Score: -199.67

Episode 200 Average Score: -197.59

Episode 300 Average Score: -198.87

Episode 400 Average Score: -200.00

Episode 500 Average Score: -200.00

Episode 600 Average Score: -200.00

Episode 700 Average Score: -200.00

Episode 800 Average Score: -200.00

Episode 900 Average Score: -200.00

Episode 1000 Average Score: -200.00

Episode 1100 Average Score: -200.00

Episode 1200 Average Score: -200.00

Episode 1300 Average Score: -200.00

Episode 1400 Average Score: -199.25

Episode 1500 Average Score: -199.41

Episode 1600 Average Score: -200.00

Episode 1700 Average Score: -200.00

Episode 1737 Average Score: -200.00