

me19b190

February 2, 2023

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
from IPython.display import display, HTML
from typing import NamedTuple, List
```

0.0.1 Gaussian Bandit Environment

```
[2]: class GaussianArm(NamedTuple):
    mean: float
    std: float

class Env:
    def __init__(self, num_arms: int, mean_reward_range: tuple, std: float):
        """
        num_arms: number of bandit arms
        mean_reward_range: mean reward of an arm should lie between
                        the given range
        std: standard deviation of the reward for each arm
        """
        self.num_arms = num_arms
        self.arms = self.create_arms(num_arms, mean_reward_range, std)

    def create_arms(self, n: int, mean_reward_range: tuple, std: float) -> dict:
        low_rwd, high_rwd = mean_reward_range
        # creates "n" number of mean reward for each arm
        means = np.random.uniform(low=low_rwd, high=high_rwd, size=(n,))
        arms = {id: GaussianArm(mu, std) for id, mu in enumerate(means)}
        return arms

    @property
    def arm_ids(self):
        return list(self.arms.keys())

    def step(self, arm_id: int) -> float:
```

```

    arm = self.arms[arm_id]
    return np.random.normal(arm.mean, arm.std)    # Reward

def get_best_arm_and_expected_reward(self):
    best_arm_id = max(self.arms, key=lambda x: self.arms[x].mean)
    return best_arm_id, self.arms[best_arm_id].mean

def get_avg_arm_reward(self):
    arm_mean_rewards = [v.mean for v in self.arms.values()]
    return np.mean(arm_mean_rewards)

def plot_arms_reward_distribution(self, num_samples=1000):
    """
    This function is only used to visualize the arm's distribution.
    """
    fig, ax = plt.subplots(1, 1, sharex=False, sharey=False, figsize=(9, 5))
    colors = sns.color_palette("hls", self.num_arms)
    for i, arm_id in enumerate(self.arm_ids):
        reward_samples = [self.step(arm_id) for _ in range(num_samples)]
        sns.histplot(reward_samples, ax=ax, stat="density", kde=True, bins=100,
        color=colors[i], label=f'arm_{arm_id}')
    ax.legend()
    plt.show()

```

0.0.2 Policy

```

[3]: class BasePolicy:
    @property
    def name(self):
        return 'base_policy'

    def reset(self):
        """
        This function resets the internal variable.
        """
        pass

    def update_arm(self, *args):
        """
        This function keep track of the estimates
        that we may want to update during training.
        """
        pass

    def select_arm(self) -> int:
        """
        It returns arm_id

```

```

"""
raise Exception("Not Implemented")

```

Random Policy

```

[4]: class RandomPolicy(BasePolicy):
    def __init__(self, arm_ids: List[int]):
        self.arm_ids = arm_ids

    @property
    def name(self):
        return 'random'

    def reset(self) -> None:
        """No use."""
        pass

    def update_arm(self, *args) -> None:
        """No use."""
        pass

    def select_arm(self) -> int:
        return np.random.choice(self.arm_ids)

```

```

[18]: class EpGreedyPolicy(BasePolicy):
    def __init__(self, epsilon: float, arm_ids: List[int]):
        self.epsilon = epsilon
        self.arm_ids = arm_ids
        self.Q = {id: 0 for id in self.arm_ids}
        self.num_pulls_per_arm = {id: 0 for id in self.arm_ids}

    @property
    def name(self):
        return f'ep-greedy ep:{self.epsilon}'

    def reset(self) -> None:
        self.Q = {id: 0 for id in self.arm_ids}
        self.num_pulls_per_arm = {id: 0 for id in self.arm_ids}

    def update_arm(self, arm_id: int, arm_reward: float) -> None:
        # your code for updating the Q values of each arm
        sum_reward = self.Q[arm_id]*self.num_pulls_per_arm[arm_id]
        sum_reward = sum_reward+arm_reward
        self.num_pulls_per_arm[arm_id] +=1
        self.Q[arm_id] = (sum_reward)/(self.num_pulls_per_arm[arm_id])

```

```

def select_arm(self) -> int:
    # your code for selecting arm based on epsilon greedy policy
    arm_maxQ = list(self.Q.keys())[list(self.Q.values()).index(max(self.Q.
↪values()))]
    probabilities = {id : 1-self.epsilon+((self.epsilon)/(len(self.arm_ids)))
↪if id == arm_maxQ else (self.epsilon)/(len(self.arm_ids)) for id in self.
↪arm_ids}
    arm_selected = np.random.choice(self.arm_ids, 1, p = list(probabilities.
↪values()))
    return int(arm_selected)

```

```

[6]: class SoftmaxPolicy(BasePolicy):
    def __init__(self, tau, arm_ids):
        self.tau = tau
        self.arm_ids = arm_ids
        self.Q = {id: 0 for id in self.arm_ids}
        self.num_pulls_per_arm = {id: 0 for id in self.arm_ids}

    @property
    def name(self):
        return f'softmax tau:{self.tau}'

    def reset(self):
        self.Q = {id: 0 for id in self.arm_ids}
        self.num_pulls_per_arm = {id: 0 for id in self.arm_ids}

    def update_arm(self, arm_id: int, arm_reward: float) -> None:
        # your code for updating the Q values of each arm
        sum_reward = self.Q[arm_id]*self.num_pulls_per_arm[arm_id]
        sum_reward = sum_reward+arm_reward
        self.num_pulls_per_arm[arm_id] +=1
        self.Q[arm_id] = (sum_reward)/(self.num_pulls_per_arm[arm_id])

    def select_arm(self) -> int:
        # your code for selecting arm based on softmax policy
        arr = np.array(list(self.Q.values()))/self.tau
        arr = np.exp(arr)
        arr = [np.exp(709) if i == np.inf else i for i in arr]
        sum_arr = sum(arr)
        probabilities = arr/sum_arr
        arm_selected = np.random.choice(self.arm_ids, 1, p = probabilities)
        return int(arm_selected)

```

```

[30]: class UCB(BasePolicy):
    def __init__(self, env, arm_ids):
        self.arm_ids = arm_ids

```

```

self.Q = {id: 0 for id in self.arm_ids}
self.num_pulls_per_arm = {id: 0 for id in self.arm_ids}

for arm_id in self.arm_ids:
    self.Q[arm_id] = env.step(arm_id)
    self.num_pulls_per_arm[arm_id] +=1

@property
def name(self):
    return 'UCB'

def reset(self):
    self.Q = {id: 0 for id in self.arm_ids}
    self.num_pulls_per_arm = {id: 0 for id in self.arm_ids}

def update_arm(self, arm_id: int, arm_reward: float) -> None:
    # your code for updating the Q values of each arm

    sum_reward = self.Q[arm_id]*self.num_pulls_per_arm[arm_id]
    sum_reward = sum_reward+arm_reward
    self.num_pulls_per_arm[arm_id] +=1
    self.Q[arm_id] = (sum_reward)/(self.num_pulls_per_arm[arm_id])

def select_arm(self) -> int:
    # your code for selecting arm based on softmax policy
    for arm_id in self.arm_ids:
        if self.num_pulls_per_arm[arm_id] == 0:
            return arm_id

    ucb = np.array([self.Q[arm_id]+((2*np.log(len(self.arm_ids))/self.
↪num_pulls_per_arm[arm_id])**0.5) for arm_id in self.arm_ids])
    arm_selected = list(self.Q.keys())[list(ucb).index(max(ucb))]
    return int(arm_selected)

```

Trainer

```

[8]: def train(env, policy: BasePolicy, timesteps):
    policy_reward = np.zeros((timesteps,))
    for t in range(timesteps):
        arm_id = policy.select_arm()
        reward = env.step(arm_id)
        policy.update_arm(arm_id, reward)
        policy_reward[t] = reward
    return policy_reward

```

```

def avg_over_runs(env, policy: BasePolicy, timesteps, num_runs):
    _, expected_max_reward = env.get_best_arm_and_expected_reward()
    policy_reward_each_run = np.zeros((num_runs, timesteps))
    for run in range(num_runs):
        policy.reset()
        policy_reward = train(env, policy, timesteps)
        policy_reward_each_run[run, :] = policy_reward

    # calculate avg policy reward from policy_reward_each_run
    avg_policy_rewards = np.array([policy_reward_each_run[:,i].mean() for i in
    ↪range(timesteps)]) # your code here (type: nd.array, shape: (timesteps,))

    total_policy_regret = sum(expected_max_reward - avg_policy_rewards) # your
    ↪code here (type: float)

    return avg_policy_rewards, total_policy_regret

```

```

[9]: def plot_reward_curve_and_print_regret(env, policies, timesteps=200,
    ↪num_runs=500):
    fig, ax = plt.subplots(1, 1, sharex=False, sharey=False, figsize=(10, 6))
    for policy in policies:
        avg_policy_rewards, total_policy_regret = avg_over_runs(env, policy,
    ↪timesteps, num_runs)
        print('regret for {}: {:.3f}'.format(policy.name, total_policy_regret))
        ax.plot(np.arange(timesteps), avg_policy_rewards, '-', label=policy.name)

    _, expected_max_reward = env.get_best_arm_and_expected_reward()
    ax.plot(np.arange(timesteps), [expected_max_reward]*timesteps, 'g-')

    avg_arm_reward = env.get_avg_arm_reward()
    ax.plot(np.arange(timesteps), [avg_arm_reward]*timesteps, 'r-')

    plt.legend(loc='lower right')
    plt.show()

```

0.0.3 Experiments

```


[10]: seed = 42
    np.random.seed(seed)

    num_arms = 5
    mean_reward_range = (-25, 25)
    std = 2.0

```

```
[11]: env = Env(num_arms, mean_reward_range, std)

env.plot_arms_reward_distribution()
```



```
[12]: best_arm, max_mean_reward = env.get_best_arm_and_expected_reward()
print(best_arm, max_mean_reward)
```

```
1 22.53571532049581
```

```
[13]: print(env.get_avg_arm_reward())
```


```
3.119254917081568
```

Please explore following values:

- Epsilon greedy: [0.001, 0.01, 0.5, 0.9]
- Softmax: [0.001, 1.0, 5.0, 50.0]


```
[14]: random_policy = RandomPolicy(env.arm_ids)
plot_reward_curve_and_print_regret(env, [random_policy], timesteps=200,
↪ num_runs=500)
```

```
regret for random: 3871.625
```



```
[19]: explore_epgreedy_epsilons = [0.001, 0.01, 0.5, 0.9]
      epgreedy_policies = [EpGreedyPolicy(ep, env.arm_ids) for ep in
                           ↪ explore_epgreedy_epsilons]
      plot_reward_curve_and_print_regret(env, epgreedy_policies, timesteps=200,
      ↪ num_runs=500)
```


```
regret for ep-greedy ep:0.001: 33.902
regret for ep-greedy ep:0.01: 84.852
regret for ep-greedy ep:0.5: 1981.669
regret for ep-greedy ep:0.9: 3491.759
```

```
[16]: explore_softmax_taus = [0.001, 1.0, 5.0, 50.0]
softmax_policies = [SoftmaxPolicy(tau, env.arm_ids) for tau in
    ↪ explore_softmax_taus]
plot_reward_curve_and_print_regret(env, softmax_policies, timesteps=200,
    ↪ num_runs=500)
```

<ipython-input-6-e56e94cd8809>:27: RuntimeWarning: overflow encountered in exp
arr = np.exp(arr)

```
regret for softmax tau:0.001: 1940.060
regret for softmax tau:1.0: 1457.339
regret for softmax tau:5.0: 400.901
regret for softmax tau:50.0: 3145.762
```



```
[31]: plot_reward_curve_and_print_regret(env, [UCB(env,env.arm_ids)], timesteps=200, num_runs=500)
```

regret for UCB: 96.566



Optional: Please explore different values of ϵ , τ and verify how does the behaviour changes.