

me19b190-tutorial-2

February 10, 2023

```
[1]: import numpy as np
import matplotlib.pyplot as plt
from typing import NamedTuple
from google.colab import output
```

```
[2]: SEED = 0

BOARD_COL = 3
BOARD_ROW = 3
BOARD_SIZE = BOARD_COL * BOARD_ROW

"""
Game board and actions are: {q, w, e, a, s, d, z, x, c}

q | w | e
--/---/--
a | s | d
--/---/--
z | x | c
"""

ACTIONS_KEY_MAP = {'q': 0, 'w': 1, 'e': 2,
                    'a': 3, 's': 4, 'd': 5,
                    'z': 6, 'x': 7, 'c': 8}
```

```
[3]: np.random.seed(SEED)
```

State Definition

```
[4]: def print_state(board, clear_output=False):
    if clear_output:
        output.clear()
    for i in range(BOARD_ROW):
        print('-----')
        out = '| '
        for j in range(BOARD_COL):
            if board[i, j] == 1:
                token = 'x'
            elif board[i, j] == -1:
```

```

        token = 'o'
    else:
        token = ' ' # empty position
    out += token + ' | '
    print(out)
    print('-----')

```

```
class State:
```

```

    def __init__(self, symbol):
        # the board is represented by an n * n array,
        # 1 represents the player who moves first,
        # -1 represents another player
        # 0 represents an empty position
        self.board = np.zeros((BOARD_ROW, BOARD_COL))
        self.symbol = symbol
        self.winner = 0
        self.end = None

```

```
@property
```

```

    def hash_value(self):
        hash = 0
        for x in np.nditer(self.board):
            hash = 3*hash + x + 1 # unique hash
        return hash

```

```

    def next(self, action: str):
        id = ACTIONS_KEY_MAP[action]
        i, j = id // BOARD_COL, id % BOARD_COL
        return self.next_by_pos(i, j)

```

```

    def next_by_pos(self, i: int, j: int):
        assert self.board[i, j] == 0
        new_state = State(-self.symbol) # another player turn
        new_state.board = np.copy(self.board)
        new_state.board[i, j] = self.symbol # current player choose to play at (i,
↪j) pos
        return new_state

```

```
@property
```

```

    def possible_actions(self):
        rev_action_map = {id: key for key, id in ACTIONS_KEY_MAP.items()}
        actions = []
        for i in range(BOARD_ROW):
            for j in range(BOARD_COL):
                if self.board[i, j] == 0:
                    actions.append(rev_action_map[BOARD_COL*i+j])

```

```

    return actions

def is_end(self):
    if self.end is not None:
        return self.end

    check = []
    # check row
    for i in range(BOARD_ROW):
        check.append(sum(self.board[i, :]))

    # check col
    for i in range(BOARD_COL):
        check.append(sum(self.board[:, i]))

    # check diagonal
    diagonal = 0; reverse_diagonal = 0
    for i in range(BOARD_ROW):
        diagonal += self.board[i, i]
        reverse_diagonal += self.board[BOARD_ROW-i-1, i]
    check.append(diagonal)
    check.append(reverse_diagonal)

    for x in check:
        if x == 3:
            self.end = True
            self.winner = 1    # player 1 wins
            return self.end
        elif x == -3:
            self.end = True
            self.winner = 2    # player 2 wins
            return self.end

    for x in np.nditer(self.board):
        if x == 0:            # play available
            self.end = False
            return self.end

    self.winner = 0          # draw
    self.end = True
    return self.end

```

Environment

```

[5]: class Env:
    def __init__(self):
        self.all_states = self.get_all_states()

```

```

self.curr_state = State(symbol=1)

def get_all_states(self):
    all_states = {} # is a dict with key as state_hash_value and value as
    ↪State object.
    def explore_all_substates(state):
        for i in range(BOARD_ROW):
            for j in range(BOARD_COL):
                if state.board[i, j] == 0:
                    next_state = state.next_by_pos(i, j)
                    if next_state.hash_value not in all_states:
                        all_states[next_state.hash_value] = next_state
                    if not next_state.is_end():
                        explore_all_substates(next_state)
    curr_state = State(symbol=1)
    all_states[curr_state.hash_value] = curr_state
    explore_all_substates(curr_state)
    return all_states

def reset(self):
    self.curr_state = State(symbol=1)
    return self.curr_state

def step(self, action):
    assert action in self.curr_state.possible_actions, f"Invalid {action} for
    ↪the current state \n{self.curr_state.print_state()}"
    next_state_hash = self.curr_state.next(action).hash_value
    next_state = self.all_states[next_state_hash]
    self.curr_state = next_state
    reward = 0
    return self.curr_state, reward

def is_end(self):
    return self.curr_state.is_end()

@property
def winner(self):
    result_id = self.curr_state.winner
    result = 'draw'
    if result_id == 1:
        result = 'player1'
    elif result_id == 2:
        result = 'player2'
    return result

```

Policy

```
[6]: class BasePolicy:
    def reset(self):
        pass

    def update_values(self, *args):
        pass

    def select_action(self, state):
        raise Exception('Not Implemented Error')
```

```
[7]: class HumanPolicy(BasePolicy):
    def __init__(self, symbol):
        self.symbol = symbol

    def select_action(self, state):
        assert state.symbol == self.symbol, f"Its not {self.symbol} symbol's turn"
        print_state(state.board, clear_output=True)
        key = input("Input your position: ")
        return key
```

```
[8]: class RandomPolicy(BasePolicy):
    def __init__(self, symbol):
        self.symbol = symbol

    def select_action(self, state):
        assert state.symbol == self.symbol, f"Its not {self.symbol} symbol's turn"
        return np.random.choice(state.possible_actions)
```

```
[9]: class ActionPlayed(NamedTuple):
    hash_value: str
    action: str

class MenacePolicy(BasePolicy):
    def __init__(self, all_states, symbol, tau=5.0):
        self.all_states = all_states
        self.symbol = symbol
        self.tau = tau

        # It store the number of stones for each action for each state
        self.state_action_value = self.initialize()
        # variable to store the history for updating the number of stones
        self.history = []

    def initialize(self):
        state_action_value = {}
        for hash_value, state in self.all_states.items():
```

```

        # initially all actions have 0 stones
        state_action_value[hash_value] = {action: 0 for action in state.
↪possible_actions}
        return state_action_value

def reset(self):
    for action_value in self.state_action_value.values():
        for action in action_value.keys():
            action_value[action] = 0

def print_updates(self, reward):
    print(f'Player with symbol {self.symbol} updates the following history with_
↪{reward} stone')
    for item in self.history:
        board = np.copy(self.all_states[item.hash_value].board)
        id = ACTIONS_KEY_MAP[item.action]
        i, j = id//BOARD_COL, id%BOARD_COL
        board[i, j] = self.symbol
        print_state(board)

def update_values(self, reward, show_update=False):
    # reward: if wins receive reward of 1 stone for the chosen action
    #         else -1 stone.
    # reward is either 1 or -1 depending upon if the player has won or lost the_
↪game.

    if show_update:
        self.print_updates(reward)
    for item in self.history:

        # your code here
        self.state_action_value[item.hash_value][item.action] = self.
↪state_action_value[item.hash_value][item.action]+reward # update_
↪state_action with appropriate term.

        self.history = []

def select_action(self, state): # Softmax action probability
    assert state.symbol == self.symbol, f"Its not {self.symbol} symbol's turn"
    action_value = self.state_action_value[state.hash_value]
    max_value = action_value[max(action_value, key=action_value.get)]
    exp_values = {action: np.exp((v-max_value) / self.tau) for action, v in_
↪action_value.items()}
    normalizer = np.sum([v for v in exp_values.values()])
    prob = {action: v/normalizer for action, v in exp_values.items()}
    action = np.random.choice(list(prob.keys()), p=list(prob.values()))
    self.history.append(ActionPlayed(state.hash_value, action))

```

```
return action
```

Game Board

```
[10]: class Game:
    def __init__(self, env, player1, player2):
        self.env = env
        self.player1 = player1
        self.player2 = player2
        self.show_updates = False

    def alternate(self):
        while True:
            yield self.player1
            yield self.player2

    def train(self, epochs=1_00_000):
        game_results = []
        player1_reward_map = {'player1': 1, 'player2': -1, 'draw': 0}
        for _ in range(epochs):
            result = self.play()

            # if player1 wins add 1 stone for the action chosen
            player1_reward = player1_reward_map[result]
            player2_reward = -player1_reward # if player2 wins add 1 stone

            self.player1.update_values(player1_reward)
            self.player2.update_values(player2_reward)

    def play(self):
        alternate = self.alternate()
        state = self.env.reset()
        while not self.env.is_end():
            player = next(alternate)
            action = player.select_action(state)
            state, _ = self.env.step(action)
        result = self.env.winner
        return result
```

Experiment

```
[11]: env = Env()

player1 = MenacePolicy(env.all_states, symbol=1)
player2 = MenacePolicy(env.all_states, symbol=-1)
# player2 = RandomPolicy(symbol=-1)
```

```
[12]: game = Game(env, player1, player2)
game.train(epochs=1_00_000)
```

```
[14]: game_with_human_player = Game(env, player1, HumanPolicy(symbol=-1))

game_with_human_player.play()

result = env.winner
print(f"winner: {result}")

player1_reward_map = {'player1': 1, 'player2': -1, 'draw': 0}
player1.update_values(player1_reward_map[result], show_update=True)
```

```
-----
| x | o | x |
-----
| o | o | x |
-----
|   | x |   |
-----
```

Input your position: c

winner: draw

Player with symbol 1 updates the following history with 0 stone

```
-----
|   |   | x |
-----
|   |   |   |
-----
|   |   |   |
-----
-----
| x |   | x |
-----
|   | o |   |
-----
|   |   |   |
-----
-----
| x | o | x |
-----
|   | o |   |
-----
|   | x |   |
-----
-----
| x | o | x |
-----
```


	o		o		x	

			x			

	x		o		x	

	o		o		x	

	x		x		o	

[13]:

orial-3-value-and-policy-iteration

February 10, 2023

```
[17]: import numpy as np
      from enum import Enum
      import copy
```

Consider a standard grid world, where only 4 (up, down, left, right) actions are allowed and the agent deterministically moves accordingly, represented as below. Here yellow is the start state and white is the goal state.

Say, we define our MDP as: - S: 121 (11 x 11) cells - A: 4 actions (up, down, left, right) - P: Deterministic transition probability - R: -1 at every step - gamma: 0.9

Our goal is to find an optimal policy (shown in right).

```
[18]: # Above grid is defined as below:
      #   - 0 denotes an navigable tile
      #   - 1 denotes an obstruction/wall
      #   - 2 denotes the start state
      #   - 3 denotes an goal state

      # Note: Here the upper left corner is defined as (0, 0)
      #         and lower right corner as (m-1, n-1)

      # Optimal Path: RIGHT RIGHT UP UP LEFT LEFT UP UP UP UP UP UP LEFT LEFT DOWN
      ↪DOWN LEFT LEFT

      GRID_WORLD = np.array([
          [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1],
          [1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1],
          [1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1],
          [1, 3, 0, 0, 1, 0, 1, 0, 1, 0, 1],
          [1, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1],
          [1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1],
          [1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1],
          [1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1],
          [1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1],
```

```

    [1, 0, 0, 0, 0, 2, 0, 0, 0, 0, 1],
    [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1]
])

```

0.0.1 Actions

```

[19]: class Actions(Enum):
    UP    = (0, (-1, 0)) # index = 0, (xaxis_move = -1 and yaxis_move = 0)
    DOWN  = (1, (1, 0))  # index = 1, (xaxis_move = 1 and yaxis_move = 0)
    LEFT  = (2, (0, -1)) # index = 2, (xaxis_move = 0 and yaxis_move = -1)
    RIGHT = (3, (0, 1))  # index = 3, (xaxis_move = 0 and yaxis_move = 1)

    def get_action_dir(self):
        _, direction = self.value
        return direction

    @property
    def index(self):
        indx, _ = self.value
        return indx

    @classmethod
    def from_index(cls, index):
        action_index_map = {a.index: a for a in cls}
        return action_index_map[index]

```

```

[20]: # How to use Action enum
for a in Actions:
    print(f"name: {a.name}, action_id: {a.index}, direction_to_move: {a.
    ↪get_action_dir()}")

print("\n-----\n")

# find action enum from index 0
a = Actions.from_index(0)
print(f"0 index action is: {a.name}")

```

```

name: UP, action_id: 0, direction_to_move: (-1, 0)
name: DOWN, action_id: 1, direction_to_move: (1, 0)
name: LEFT, action_id: 2, direction_to_move: (0, -1)
name: RIGHT, action_id: 3, direction_to_move: (0, 1)

```

```

-----
0 index action is: UP

```

0.0.2 Policy

```
[21]: class BasePolicy:
    def update(self, *args):
        pass

    def select_action(self, state_id: int) -> int:
        raise NotImplemented

class DeterministicPolicy(BasePolicy):
    def __init__(self, actions: np.ndarray):
        # actions: its a 1d array (|S| size) which contains action for each state
        self.actions = actions

    def update(self, state_id, action_id):
        assert state_id < len(self.actions), f"Invalid state_id {state_id}"
        assert action_id < len(self.actions), f"Invalid action_id {action_id}"
        self.actions[state_id] = action_id

    def select_action(self, state_id: int) -> int:
        assert state_id < len(self.actions), f"Invalid state_id {state_id}"
        return self.actions[state_id]
```

0.0.3 Environment

```
[22]: class Environment:
    def __init__(self, grid):
        self.grid = grid
        m, n = grid.shape
        self.num_states = m*n

    def xy_to_posid(self, x: int, y: int):
        _, n = self.grid.shape
        return x*n + y

    def posid_to_xy(self, posid: int):
        _, n = self.grid.shape
        return (posid // n, posid % n)

    def isvalid_move(self, x: int, y: int):
        m, n = self.grid.shape
        return (x >= 0) and (y >= 0) and (x < m) and (y < n) and (self.grid[x, y] != 1)

    def find_start_xy(self) -> int:
        m, n = self.grid.shape
```

```

for x in range(m):
    for y in range(n):
        if self.grid[x, y] == 2:
            return (x, y)
raise Exception("Start position not found.")

def find_path(self, policy: BasePolicy) -> str:
    max_steps = 50
    steps = 0

    P, R = self.get_transition_prob_and_expected_reward()
    num_actions, num_states = R.shape
    all_possible_state_posids = np.arange(num_states)

    path = ""
    curr_x, curr_y = self.find_start_xy()
    while (self.grid[curr_x, curr_y] != 3) and (steps < max_steps):
        curr_posid = self.xy_to_posid(curr_x, curr_y)
        action_id = policy.select_action(curr_posid)
        next_posid = np.random.choice(
            all_possible_state_posids, p=P[action_id, curr_posid])
        action = Actions.from_index(action_id)
        path += f" {action.name}"
        curr_x, curr_y = self.posid_to_xy(next_posid)
        steps += 1
    return path

def get_transition_prob_and_expected_reward(self): #  $P(s_{next} | s, a)$ ,  $R(s, a)$ 
    m, n = self.grid.shape
    num_states = m*n
    num_actions = len(Actions)
    P = np.zeros((num_actions, num_states, num_states))
    R = np.zeros((num_actions, num_states))
    for a in Actions:
        for x in range(m):
            for y in range(n):
                xmove_dir, ymove_dir = a.get_action_dir()
                xnew, ynew = x + xmove_dir, y + ymove_dir # find the new co-ordinate
                # after the action a

                posid = self.xy_to_posid(x, y)
                new_posid = self.xy_to_posid(xnew, ynew)

                if self.grid[x, y] == 3:
                    # the current state is a goal state

```

```

        P[a.index, posid, posid] = 1
        R[a.index, posid] = 0
    elif (self.grid[x, y] == 1) or (not self.isvalid_move(xnew, ynew)):
        # the current state is a block state or the next state is invalid
        P[a.index, posid, posid] = 1
        R[a.index, posid] = -1
    else:
        # action a is valid and goes to a new position
        P[a.index, posid, new_posid] = 1
        R[a.index, posid] = -1
    return P, R

```

0.0.4 Policy Iteration

```

[23]: def policy_evaluation(P: np.ndarray, R: np.ndarray, gamma: float,
                             policy: BasePolicy, theta: float,
                             init_V: np.ndarray=None):
    num_actions, num_states = R.shape

    # Please try different starting point for V you will find it will always
    # converge to the same V_pi value.
    if init_V is None:
        init_V = np.zeros(num_states)
    V = copy.deepcopy(init_V)

    delta = 100.0
    while delta > theta:
        delta = 0.0
        for state_id in range(num_states):
            action_id = policy.select_action(state_id)
            v_old = V[state_id]
            # Following equation is a different way of writing the same equation
            → given in the slide.
            # Note here R is an expected reward term.
            V[state_id] = R[action_id, state_id] + gamma * np.dot(P[action_id,
            → state_id], V)
            delta = max(delta, abs(V[state_id] - v_old))
    return V

def policy_improvement(P: np.ndarray, R: np.ndarray, gamma: float,
                       policy: BasePolicy, V: np.ndarray):
    num_actions, num_states = R.shape
    policy_stable = True
    for state_id in range(num_states):
        old_action_id = policy.select_action(state_id)

```

```

    # your code here
    new_action = old_action_id
    V_old = V[state_id]
    for action in Actions:
        action_id = action.index
        q_new = R[action_id, state_id] + gamma * np.dot(P[action_id, state_id], V)
        if q_new > V_old:
            new_action = action_id

    new_action_id = new_action # update new_action_id based on the value
    ↪function.

    policy.update(state_id, new_action_id)
    if old_action_id != new_action_id:
        policy_stable = False
    return policy_stable

def policy_iteration(P: np.ndarray, R: np.ndarray, gamma: float,
                    theta: float=1e-3, init_policy: BasePolicy = None):
    num_actions, num_states = R.shape

    # Please try exploring different policies you will find it will always
    # converge to the same optimal policy for valid states.
    if init_policy is None:
        # Say initial policy = all up actions.
        init_policy = DeterministicPolicy(actions=np.zeros(num_states, dtype=int))

    # creating a copy of a initial policy
    policy = copy.deepcopy(init_policy)
    policy_stable = False
    while not policy_stable:
        V = policy_evaluation(P, R, gamma, policy, theta)
        policy_stable = policy_improvement(P, R, gamma, policy, V)
    return policy, V

```

0.0.5 Value Iteration

```

[24]: # Directly find the optimal value function
def get_optimal_value(P: np.ndarray, R: np.ndarray, gamma: float,
                    theta: float, init_V: np.ndarray=None):
    num_actions, num_states = R.shape

    # Please try different starting point for V you will find it will always
    # converge to the same V_star value.
    if init_V is None:
        init_V = np.zeros(num_states)

```

```

V = copy.deepcopy(init_V)

delta = 100.0
while delta > theta:
    delta = 0.0
    for state_id in range(num_states):
        v_old = V[state_id]
        q_sa = np.zeros(num_actions)
        for a in Actions:
            q_sa[a.index] = R[a.index, state_id] + gamma * np.dot(P[a.index,
↪state_id], V)
        V[state_id] = np.max(q_sa)
        delta = max(delta, abs(V[state_id] - v_old))
    return V

def value_iteration(P: np.ndarray, R: np.ndarray, gamma: float,
                    theta: float=1e-3, init_V: np.ndarray=None):
    V_star = get_optimal_value(P, R, gamma, theta, init_V)

    num_actions, num_states = R.shape
    policy = DeterministicPolicy(actions=np.zeros(num_states, dtype=int))
    for state_id in range(num_states):
        # Your code here

        best_action = policy.select_action(state_id)
        state_value_after_best_action = -np.inf
        n = P.shape[-1]
        n = int(n**0.5)
        temp_grid = np.empty((n,n))
        for action in Actions:
            x_curr, y_curr = env.posid_to_xy(state_id)
            x_dir, y_dir = action.get_action_dir()
            x_new, y_new = (x_curr+x_dir) , (y_curr+y_dir)
            if env.isvalid_move(x_new, y_new):
                state_transitioned = env.xy_to_posid(x_new, y_new)
                V_transition = V_star[state_transitioned]
                if V_transition > state_value_after_best_action: #greedy with respect
↪to optimal value function
                    best_action = action.index
                    state_value_after_best_action = V_transition

        action_id = best_action # update the action_id based on V_star

    policy.update(state_id, action_id)

    return policy, V_star

```


0.0.6 Experiments

```
[25]: def is_same_optimal_value(V1, V2, diff_theta=1e-3):  
      diff = np.abs(V1 - V2)  
      return np.all(diff < diff_theta)
```

```
[26]: seed = 0  
      np.random.seed(seed)  
  
      gamma = 0.9  
      theta = 1e-5
```

```
[27]: env = Environment(GRID_WORLD)  
      P, R = env.get_transition_prob_and_expected_reward()
```

Exp 1: Using Policy iteration algorithm find the optimal path from start to goal position

```
[28]: # # Start with random choice of init_policy.  
      # One such choice could be: init_policy = np.ones(env.num_states, dtype=int)  
      # Start with your own choice of init_policy  
      init_policy = DeterministicPolicy(actions=np.ones(env.num_states, dtype=int))  
  
      pitr_policy, pitr_V_star = policy_iteration(P, R, gamma, theta=theta,  
      ↪ init_policy=init_policy)  
      pitr_path = env.find_path(pitr_policy)  
      print(pitr_path)
```

RIGHT RIGHT UP UP LEFT LEFT UP UP UP UP UP UP LEFT LEFT DOWN DOWN LEFT LEFT

Exp 2: Using value iteration algorithm find the optimal path from start to goal position

```
[29]: vitr_policy, vitr_V_star = value_iteration(P, R, gamma, theta=theta)  
      vitr_path = env.find_path(vitr_policy)  
      print(vitr_path)
```

RIGHT RIGHT UP UP LEFT LEFT UP UP UP UP UP UP LEFT LEFT DOWN DOWN LEFT LEFT

Exp 3: Compare the optimal value function of policy iteration and value iteration algorithm

```
[30]: is_same_optimal_value(pitr_V_star, vitr_V_star)
```

```
[30]: True
```

Exp 4: Using initial guess for V as random values, find the optimal value function using policy evaluation and compare it with the optimal value function

```
[31]: # Start with random choice of init_V.
# One such choice could be: init_V = np.random.randn(env.num_states)
# Another choice could be: init_V = 10*np.ones(env.num_states)
# Start with your own choice of init_V
init_V = 10*np.ones(env.num_states) # your choice

V_star = policy_evaluation(P, R, gamma, pitr_policy, theta, init_V)
is_same_optimal_value(pitr_V_star, V_star)
```

[31]: True

Exp 5: Using initial guess for V as random values, find the optimal value function using `get_optimal_value` and compare it with the optimal value function

```
[32]: # Start with random choice.
# One such choice could be: init_V = np.random.randn(env.num_states)
# Another choice could be: init_V = 10*np.ones(env.num_states)
# Start with your own choice of init_V
init_V = 100*np.ones(env.num_states)

V_star = get_optimal_value(P, R, gamma, theta, init_V)
is_same_optimal_value(vitr_V_star, V_star)
```

[32]: True

[32]:

Exp Optional: Try changing the grid by adding multiple paths to the goal state and check if our `policy_iteration` or `value_iteration` algorithm is able to find optimal path. Redo the above experiments.

- 1 way to add another path would be `GRID_WORLD[4, 1] = 0`

[32]: